

Contents

Conference-paper corpus study, 11 six-page papers against our ICCIT draft	1
1. The corpus roster	1
2. Paper-level metrics, ours against the eleven-paper band	1
3. Per-section word budget	2
4. Structural conventions	2
5. Sentence-length distribution	2
6. Read before quoting any number above	3

Conference-paper corpus study, 11 six-page papers against our ICCIT draft

How to read the columns. *ours* is `paper/iccit6/main.pdf`, our six-page ICCIT submission. *min*, *median*, *max* are the band over the eleven reference papers; our paper is never a member of its own band, so every band here is $n=11$. *pct* is our percentile inside that band, printed on every row because a min..max band over eleven papers passes at the extreme. A flag of OUT-HIGH or OUT-LOW means outside the band entirely.

What the corpus is, and what it is not. Eleven papers of six pages or fewer, pulled from four of our other project reference folders (SAR drone, ML paper, XAI NIDS). They are not a sample of ICCIT and not a sample of any single venue: ten of the eleven were read from arXiv-hosted copies, and only one (Okabe, Interspeech 2018) carries a printed conference venue line. Topic spread is wide. Read a band as “the norm among these eleven short papers”, never as “the ICCIT norm”.

What was changed in the instrument. The measurement code is `scripts/78_corpus_style_study.py` and `src/markdown_corpus.py`, copied from the ML_Paper repository, where it was written for twelve-to-twenty page NDSS papers. Four changes were made and no measurement rule beyond them. Three are recalibrations for six-page papers, namely the body-word sanity window from 5,000..20,000 down to 1,500..20,000, the minimum detected sections from five to four, and the references-heading pattern now also accepts arabic numbering (6. References, the Interspeech style). Image markdown is stripped before measurement because this corpus was converted with figure extraction on, unlike the NDSS bundle. The fourth is a parsing fix found by reading the output rather than the page count, since the caption pattern allowed one page-anchor span before the label where marker sometimes emits two, which silently cost a figure. It now accepts any number of anchors, and every count in this report, corpus and ours alike, was recomputed under that rule.

1. The corpus roster

paper	pp	body words	sent	sec	tab	fig	med w/sent
OURS 00_OURS_iccit6	6	3642	179	6	5	2	20
AlvarezMelis2018_robustness-interpretability	6	2464	126	4	0	8	18.0
Anon2025_uav-vlrr-nmpc	6	2654	140	8	3	7	19.0
Batool2026_humandiffusion-sar	5	1873	102	6	0	4	18.0
Boddu2025_quantized-yolov4tiny-pi	6	2874	158	6	1	4	18.0
Manzini2023b_differentiable-boustrophedon	6	3671	181	9	1	5	18
Yang2022_rgb-ir-fusion-drone	6	3341	137	6	6	7	21
ref17_gueriani_resnet_bigru_arxiv2026	6	2744	146	6	3	4	18.0
ref18_gueriani_vit_bilstm_arxiv2026	6	2292	113	5	4	4	20
ref28_okabe_attentivestats_interspeech2018	5	2563	128	5	3	2	17.5
ref32_yang_lnetskd_arxiv2023	6	2930	158	5	4	4	18.0
ref34_debicha_advtrain_arxiv2021	5	3022	130	5	1	4	20.0

Body words means first numbered heading to REFERENCES. Front matter and the bibliography are outside it, so a body-word figure is not a page count in words.

2. Paper-level metrics, ours against the eleven-paper band

metric	ours	min	median	max	pct	flag
body words	3642	1873	2744	3671	91	high
sentences	179	102	137	181	91	high
mean words / sentence	20.3	18.2	19.6	24.4	82	-
median words / sentence	20	17.5	18.0	21	91	high
% sentences over 35w	5.0	1.9	8.0	16.2	36	-
longest sentence	59	38	53	103	64	-
paragraphs	63	37	49	94	73	-
mean words / paragraph	57.8	30.6	50.6	90.3	82	-
citations per 1k	11.8	0.0	10.6	18.2	73	-
first person per 1k	4.9	0.0	3.0	16.2	55	-
passive per 1k [proxy]	9.1	8.5	15.7	30.0	9	low
% sentences past [proxy]	5.6	0.6	12.3	30.4	27	-
% no tense anchor	38.0	17.7	38.4	54.0	36	-
sections	6	4	6	9	82	-
sections (numbered)	6	4	5	8	82	-
tables (total)	5	0	3	6	91	high
figures (total)	2	2	4	8	9	low

3. Per-section word budget

role	n papers	ours	min	median	max	verdict
introduction	11	534	301	493	671	in band, median 493
background	1	ABSENT	818	818	818	no band, only 1 of 11 papers
related_work	5	548	342	372	725	in band, median 372
threat_model	2	ABSENT	287	822	1357	no band, only 2 of 11 papers
method	4	849	532	743	973	in band, median 743
evaluation	10	1131	462	874	1575	in band, median 874
discussion	2	460	405	503	601	no band, only 2 of 11 papers, ours inside the pair
conclusion	9	120	74	121	325	in band, median 121

A role carried by fewer than three of the eleven papers has no band. Background appears in one reference paper and a threat model in two, so neither is a convention this corpus can be said to have, and their absence from our draft is not a violation. Every role that can be banded, and every role our draft carries, is inside its band.

The conclusion is now its own section. The draft previously ran one Discussion and Conclusion section, so every word of it was filed under discussion and the conclusion role read ABSENT against nine of eleven reference papers that carry a standalone one. Splitting the two puts both inside their bands and raises the section count to six, which is the corpus median.

Roles are classified from section titles, by the same regexes for every paper, ours included. A paper whose method section is titled after its artifact rather than after its function classifies elsewhere or nowhere, which is why the n column varies by role and why an ABSENT is a statement about titles, not about content.

4. Structural conventions

convention	corpus	ours
Related Work as its own section	5/11	yes
roadmap paragraph in the introduction	2/11	yes
any back matter before REFERENCES	5/11	no
marker page anchors present	9/11	yes

5. Sentence-length distribution

paper	1-10	11-20	21-30	31-40	41-50	51+
OURS 00_OURS_iccit6	22%	32%	28%	15%	2%	1%
AlvarezMelis2018_robustness-interpretabi	19%	43%	21%	13%	2%	2%
Anon2025_uav-vlrr-nmpc	19%	41%	30%	8%	1%	1%
Batool2026_humandiffusion-sar	20%	40%	31%	8%	1%	0%
Boddu2025_quantized-yolov4tiny-pi	22%	40%	28%	10%	0%	0%
Manzini2023b_differentiable-boustrophedo	12%	49%	24%	12%	3%	0%
Yang2022_rgb-ir-fusion-drone	10%	37%	31%	11%	3%	8%
ref17_gueriani_resnet_bigru_arxiv2026	14%	47%	34%	4%	1%	1%
ref18_gueriani_vit_bilstm_arxiv2026	18%	35%	35%	8%	4%	1%
ref28_okabe_attentivestats_interspeech20	20%	42%	22%	10%	5%	2%
ref32_yang_lnetskd_arxiv2023	15%	45%	32%	6%	1%	0%
ref34_debicha_advtrain_arxiv2021	11%	42%	29%	7%	5%	6%

Percentages of that paper’s own sentences. The short-sentence column is where our draft separates from the corpus, and it is the same fact the mean and median words-per-sentence rows in Section 2 report.

6. Read before quoting any number above

Figure and table counts are caption counts. They come from markdown produced by marker, so a caption lost in conversion is an undercount with no error raised. Numeral gaps are the only detector and are reported per paper in the manifest.

The tense and passive columns are regex proxies, not parses. No part-of-speech tagger was used. Past passive increments both the past and the passive counter, so the two are not independent and must never be summed. Between 18 and 54 percent of sentences carry no tense anchor at all in this corpus; read every tense figure against that share.

Calibration was not done. The proxies have not been hand-validated against labelled sentences here any more than they were in the NDSS study.

Provenance. Manifest `report/corpus_style_study.json`, normalisation v1, 17 ordered cleaning steps. Every source markdown file is sha256-pinned in the manifest, as is the PDF it was converted from. Regenerate with `python3 scripts/corpus_style_study.py --md-dir ../md --write-manifest` then `python3 scripts/build_report.py`.